Table 1. Task success and inference cost at two token budgets. Adaptive routing is best on every suite within each budget: 2.9 and 2.4 points above Verifier, the strongest baseline, on average, with 13% and 18% fewer tokens.

Method	Task success (%) ↑					Cost ↓	
	Math n=250	Code n=200	Docs n=150	Tools n=100	Avg. macro	Tokens mean	Latency ms
16k-token budget							
Direct	65.6	61.5	73.3	54.0	63.6	**4,800**	**410**
Chain of thought	72.0	68.0	77.3	61.0	69.6	6,200	521
Search	78.8	74.0	82.0	65.0	75.0	10,150	811
Verifier	78.0	75.5	83.3	64.0	75.2	11,280	953
Adaptive routing	**80.4**	**77.5**	**85.3**	**69.0**	**78.1**	9,800	754
Δ vs. Verifier	+2.4	+2.0	+2.0	+5.0	+2.9	−1,480	−198
32k-token budget							
Chain of thought	73.6	70.5	80.7	64.0	72.2	**9,700**	**743**
Search	82.8	77.5	87.3	71.0	79.7	17,200	1,236
Verifier	83.2	79.0	88.7	72.0	80.7	18,400	1,403
Adaptive routing	**84.8**	**81.5**	**90.0**	**76.0**	**83.1**	15,100	1,128
Δ vs. Verifier	+1.6	+2.5	+1.3	+4.0	+2.4	−3,300	−274

Synthetic data. Success = successes / n; Avg. is the unweighted mean of the four rates, and orders the baselines. Δ is Adaptive routing minus Verifier, from unrounded values; green marks an improvement. Bold: best within a budget. Shaded: proposed method.