Table 2. Ablation study on PaperBananaBench. Each row removes or weakens agents of the default pipeline ①. Every change lowers the overall win rate, by 8.4 to 16.0 points.

	Agents					Win rate vs. human (%) ↑					
#	Retriever	Planner	Stylist	Visualizer	Critic	Faith.	Conc.	Read.	Aesth.	Overall	Δ Overall
①	✓	✓	✓	✓	3 iters	**45.8**	**80.7**	**51.4**	**72.1**	**60.2**	–
②	✓	✓	✓	✓	1 iter	38.3	75.2	50.6	<u>68.9</u>	<u>51.8</u>	−8.4
③	✓	✓	✓	✓	none	30.7	<u>79.2</u>	47.0	**72.1**	45.6	−14.6
④	✓	✓	✗	✓	none	39.2	61.7	47.9	67.4	49.2	−11.0
⑤	~	✓	✗	✓	none	37.3	62.7	<u>51.1</u>	65.6	48.3	−11.9
⑥	✗	✓	✗	✓	none	<u>41.9</u>	58.6	43.1	62.9	44.2	−16.0

Faith., Conc., Read., Aesth.: faithfulness, conciseness, readability, aesthetics. Critic: refinement iterations. Δ Overall: the row minus row ①, in points (computed here). Bold: best and underline: second best in each column; ties share the mark. Shaded: the default PaperBanana setting. – = not applicable. ✓ agent used, ✗ agent removed, ~ random retriever (10 random examples).