Table 4. Datasets and evaluation splits. Counts are in each collection's own unit, so no total is given.

		Split size				
Collection	**Coverage**	**Train**	**Dev**	**Test**	**Unit**	**Separation rule**
			Training and held-out			
Mathematical proofs	English	120,000	2,000	3,000	problem	Hold out entire problem families
Repository patches	Python / C++	8,200	600	900	task	Disjoint repositories; freeze revision
Tool trajectories	6 languages	45,000	1,500	2,500	trajectory	Hold out goal templates and tool schemas
Long-document QA	2 domains	18,000	800	1,200	question	Split by source document before sampling
Technical manuals	12 domains	125,000	0	5,000	document	Zero dev entries are intentional
			Evaluation only			
Sealed code tasks	Python	*NA*	*NR*	750	problem	Evaluation only; dev count unspecified
Interactive tasks	Browser / CLI	*NA*	100	400	episode	Freeze environment and initial state

Synthetic data. Counts are manifest entries; 0 is an intentionally empty split. NA = not applicable; NR = not reported.